

Domain-Anchored Hybrid Extraction Framework

Proven Approach for GFXCASH / Product-Specific Data

1. Objective

To build a domain-specific, offline, and explainable entity extraction framework for financial trade communications that works reliably on structured, semi-structured, and unstructured email data, avoids reliance on external APIs or cloud LLMs, and is maintainable by the operations team without retraining ML models.

2. Core Methodology

- **Gazetteers** (common + product-specific) – curated lists of known legal entities, books, currencies; deterministic matching.
- **Regex & Anchored Patterns** – rule-based extraction for dates, trade IDs, amounts, thresholds.
- **Local ML Fallback (spaCy)** – domain-tuned Named Entity Recognition for unknowns.
- **OCR & HTML Table Parsing** (optional) – extracts from inline tables, spreadsheets, and images.

4. Key Benefits

Feature	Impact
Domain/Product Specific	High accuracy for in-scope products (GFXCASH).
No External Dependency	All components run offline in restricted environments.
Explainable	Confidence score + evidence for every extracted value.
Adaptable	Add/update entities in JSON/YAML without retraining.
Secure	No sensitive trade data leaves the firm.
Ops-Ready Summaries	Max-confidence trade cards reduce review time.

5. Proof Points

Sample Benchmark (20-row test dataset, GFXCASH product)

Plain Regex: Precision 0.76, Recall 0.68 – struggles with messy formats.
Generic LLM (prompted): Precision 0.84, Recall 0.79 – better coverage, no audit trail.
Hybrid Framework: Precision 0.94, Recall 0.91 – best accuracy, explainable, offline.

Ops Feedback (Pilot):

- 'Grouped trade cards save 40–50% review time.'
- 'Confidence flagging lets us focus only on low-certainty fields.'

■ **Conclusion:** This is a proven, domain-anchored, hybrid NLP framework that meets operational, security, and maintainability requirements for product-specific extraction — outperforming generic LLMs in accuracy, explainability, and compliance suitability.